

Hybrid Telco Churn and Geospatial Risk Analysis (Final Report)

Executive Summary

The Problem

Traditional churn prediction in the telecommunications sector relies heavily on billing and demographic data, often overlooking the physical reality of the customer experience. This project introduces a **Hybrid Geospatial Risk Model** that bridges this "data blind spot" by proving that while billing identifies *who* is at risk, infrastructure failures—specifically a **0.6 signal stability threshold**—act as the actual "breaking point" for high-value users.

The Hybrid Approach

By scientifically comparing multiple machine learning architectures and optimizing a Logistic Regression Model through GridSearch, I achieved a robust **Test ROC-AUC of 0.84**. To address the business requirement of balancing retention costs with churn prevention, I intentionally tuned the model's decision threshold to 0.43. This "defensive strategy" prioritizes a high **Recall of 85%**, ensuring that nearly all at-risk customers in technical "dead zones" are identified for proactive intervention. While this increases false alarms, the cost of a remote signal check is negligible compared to the \$500+ cost of losing a subscriber.

Key Findings

- **Fiber Risk Proxy:** Feature importance analysis identifies *Internet Service: Fiber optic* as the #1 most critical predictor of churn. Because Fiber customers have the highest performance expectations, they are the most sensitive to signal degradation, making "Fiber Optic" a major proxy for geospatial risk.
- **The Tipping Point:** Our EDA confirms that infrastructure instability acts as a "trigger event"; high-value customers are significantly more likely to leave when reliability drops below the identified 0.6 stability threshold, regardless of their global ranking.
- **Dominance of Financial “Anchors”:** Traditional metrics like *Monthly Charges* and *Contract: Two-year* remain top drivers (#2 and #3), confirming that while technical issues

cause the frustration, financial commitment is what currently keeps customers from leaving immediately.

Strategic Recommendations

I recommend shifting from reactive discounting to a **Geospatial Retention Strategy**. This includes prioritizing technical maintenance in identified "Churn Hotspots" and utilizing the model's high-recall output to incentivize at-risk users into **One-Year contracts** to anchor behavior while infrastructure repairs are performed.

1.0 Introduction and Objectives

In today's hyper-connected world, the telecommunications industry plays a critical role in facilitating communication at both local and global levels. Companies in this sector face relentless competition, making customer loyalty vital to sustaining profitability. The loss of a single customer not only reduces immediate revenue but also allows competitors to capture market share, underscoring the importance of churn prediction for proactive business strategies.

Many researchers have traditionally relied on churn classification methods to analyze billing information, service types, and transaction data to predict customer behavior. However, these methods often overlook the impact of physical infrastructure on churn dynamics, such as service reliability and network stability.

This context necessitates a more comprehensive approach, the Hybrid Geospatial Risk model, which integrates traditional billing data with critical metrics related to network infrastructure. By evaluating parameters like signal stability index and outage history score. This project introduces a Hybrid Geospatial Risk model, which augments traditional billing data with engineered metrics related to network infrastructure, such as the signal stability index and outage history score. The ultimate objective is to synthesize these diverse perspectives into an optimized predictive framework, balanced for both high recall and precision, that allows telecom companies to allocate resources effectively through proactive, localized retention strategies.

2.0 Data Sourcing and Cleaning

2.0.1 Dataset Description

The dataset utilized for this analysis is the Telco Customer Churn dataset (IBM Sample Data), which contains information on 7,043 rows (customers) and 21 columns (initial features). Each row represents a unique customer, and the columns provide a mix of demographic, service-related, and financial data. The 21 columns are categorized as follows:

Demographics:

- Gender, senior citizen (binary), partner, and dependents (whether the customer has a spouse or children).

Customer Account Information:

- Tenure (number of months the customer has stayed with the company).
- Contract (month-to-month, one-year, two-year).
- Paperless billing and payment method (electronic check, mailed check, bank transfer, credit card).

Service Information:

- Phone service and multiple lines.
- Internet Service (DSL, Fiber optic, No).
- Online add-ons: online security, online backup, device protection, tech support, streaming TV, and streaming movies

Financial Metrics:

- Monthly charges and total charges.

Target Variable:

- Churn (Whether the customer left within the last month).

To ensure data integrity, I performed the following cleaning steps:

- **Missing Values and Duplicates:** An initial check (using `.isnull().sum()` and `.duplicated().sum()`) indicated zero null values and no duplicate records.
- **Handling Empty Strings (" "):** Despite the initial null check, a closer examination of the total charges column revealed empty strings (" ") for 11 new customers with zero months of tenure. I utilized `pd.to_numeric` with `errors='coerce'` to convert these artifacts into a numeric format and subsequently filled them with 0, representing a \$0 baseline for customers in their first billing cycle. This ensured the column was correctly formatted as a float for mathematical analysis.
- **Target Binarization:** The target variable, Churn, was converted from text ("Yes"/"No") into a binary format (1 for 'Yes', 0 for 'No') to ensure the target was isolated from the predictor set, preventing data leakage.
- **Feature Selection:** The customer ID column was removed to prevent the model from "memorizing" specific individuals. Furthermore, to ensure financial metrics and tenure did not disproportionately influence the model due to their large numerical ranges, I applied **StandardScaler** to all continuous variables. To prevent data leakage, the scaler was fitted exclusively on the training sets, ensuring the model prioritized patterns over raw magnitudes while maintaining strict separation of unseen data

This meticulous cleaning process ensured 100% data density, providing a stable foundation for the modeling phase without risk of technical errors or biased results.

3.0 Hybrid Feature Injection and Engineering

The core innovation of this case study is the hybrid modeling approach, which combines traditional behavioral data with engineered infrastructure metrics. By simulating regional service differences across five specific Zip codes, I integrated the physical environment into the predictive engine.

Customers were assigned to specific neighborhoods to represent regional infrastructure variance. For example, Zip 98104 was designated as a 'Trouble Zone' with poor tower density, allowing me to test the impact of technical failure on customer loyalty.

I engineered three key geospatial features to quantify infrastructure health:

- **Signal Stability Index:** a localized measure of network reliability.
- **Outage History Score:** the frequency of service interruptions in a region.
- **Competitor Exposure:** a binary indicator of whether rival service providers target a high-risk zone.

After injecting these three new metrics and incorporating the Zip Code, the feature count briefly increased to 24 columns (excluding customer ID). This prevents the model from ‘cheating’ by memorizing the exact regional average and forces it to learn the underlying relationship between signal degradation and churn. This process resulted in a high-dimensional feature set of over 40 numerical and binary predictors post One-Hot Encoding for categorical variables (like Contract and Internet Service). Furthermore, to ensure model robustness, I applied Gaussian noise to the signal stability index, better simulating real-world signal fluctuations and preventing overfitting on static regional data.

4.0 Exploratory Data Analysis (EDA)

In this phase, I conducted a dual analysis of customer behavioral metrics and geographical infrastructure factors. This comprehensive view allows us to identify key drivers of churn and inform strategies that telecommunications companies can adopt to improve customer retention.

4.0.1 Tenure Distribution and Churn Propensity

Before layering in geospatial risk, it is essential to establish the behavioral baseline of the dataset. This analysis examines Tenure against Churn to determine the 'Critical Window' where customers are most likely to exit the service.




Figure 1: Tenure Distribution by Churn Status

The boxplot reveals a stark contrast in customer loyalty. The median tenure for retained customers (No) sits near 40 months, whereas the median for churned customers (Yes) is significantly lower, at approximately 10 months. Notably, 75% of all churn occurs within the first 30 months of service.

For our Geospatial Risk Analysis, this identifies a high-priority 'vulnerability window.' If a new customer (Tenure < 20 months) is in a high-risk Zip Code with poor Signal Stability, they represent the highest possible churn risk to the organization.

The presence of high-tenure outliers in the churned group suggests that even long-term 'loyal' customers are susceptible to exit, likely driven by the localized infrastructure failures investigated in the hybrid model.

4.0.2 Total Charges vs. Churn in High-Risk Zip Codes

In this multivariate analysis, I examine the relationships among location (Zip Code), customer financial commitment (Total Charges), and Churn. By isolating the top 4 high-risk zip codes, we can see if geospatial location acts as a key variable in the churn equation.




Figure 2: Total Charges vs. Churn in High-Risk Zip Codes

Across all four zip codes (98101, 98102, 98103, 98105), we observe a systemic risk pattern. The teal 'Yes' (Churn) peak is significantly higher than the pink 'No' peak at the low end of the spectrum (near \$0 Total Charges). This confirms that in these specific geospatial areas, customers are departing early in their lifecycle. The high density of churn at low total charges suggests that **regional infrastructure issues may be preventing customers from reaching service maturity**, creating a 'revolving door' effect in these high-risk zones.

4.0.3 Churn Rate Sensitivity to Signal Stability

This analysis explores the direct relationship between network reliability (Signal Stability Index) and Churn Rate. By binning the stability scores, we can pinpoint the exact 'breaking point' where technical frustration leads to a significant spike in customer departure.




Figure 3: Churn Rate by Signal Stability Index

This visualization reveals a dramatic surge in churn when signal stability hits the 0.6 mark, where the churn rate at 0.6 is roughly **5-10% higher** than at more stable levels (0.7-1.0). While customers seem to tolerate moderate instability (0.7 to 0.9) with a relatively flat churn rate, **the 0.6 threshold** represents a critical **technical breaking point**. Interestingly, the churn rate at 'Perfect Signal' (1.0) remains comparable to higher stability tiers, confirming that once technical reliability is achieved, behavioral and financial metrics resume their role as the primary determinants of churn.

4.0.4 Total Charges as a Financial Risk Profile

In this analysis, I use total charges to understand the spending profile of customers who churn. This distribution helps identify a critical 'early-exit' window because total charges are a direct reflection of tenure.




Figure 4: Density Profile of Total Charges and Churn

The massive teal peak at low Total Charges (under \$1,000) confirms a high-risk segment of customers who leave almost immediately. While this specific chart focuses on spending, it points toward a 'Perfect Storm' of risk: likely new customers on Month-to-Month contracts who encounter early service friction. Their low total spending indicates they are exiting before the service can prove its long-term value. Notably, the **retention probability begins to exceed churn risk once customers surpass the \$2,000 spending threshold**, identifying this as the critical 'anchor point' for long-term loyalty

4.0.5 The Technical Breaking Point of Loyalty

This analysis explores the interaction between Contract Type, Signal Stability, and Churn. By establishing a 'Trouble Zone Threshold' at 0.6, we can identify exactly when technical instability becomes the primary driver for customer departure, regardless of their contract terms.




Figure 5: Signal Stability Impact across Contract Types

While most customers across all contract types maintain stability above the 0.6 threshold, the Two-year contract group reveals a critical finding. Although these customers are typically the most 'locked-in,' we see a cluster of churn outliers deep within the 'Trouble Zone' (below 0.5). This suggests that while long-term contracts provide a financial safety net, extreme signal instability acts as a **'contract breaker.'** Once service quality drops deep into the Trouble Zone, even the most loyal, long-term customers will exit the service, proving that **technical reliability is the ultimate prerequisite for retention.**

5.0 Data Preprocessing and Transformation

Following the insights from the EDA, the dataset was prepared for the machine learning phase through the following key structural changes:

5.0.1 Categorical Encoding (One-Hot Encoding):

Since machine learning models require numerical input, I converted all text-based categories (such as payment method, contract, and Internet Service) into a binary format.

I utilized `pd.get_dummies()` with the `drop_first=True` parameter. Dropping the first column of each category prevents the "Dummy Variable Trap" (multicollinearity), ensuring that the model does not receive redundant information that could confuse its logic. This process expanded the feature set from 24 columns to over 40 predictors. Despite the increased dimensionality, data integrity was maintained, as each new feature serves as a clear binary indicator (0 or 1).

5.0.2 Numerical Scaling and Normalization

With the categorical data encoded, I addressed the variance in the continuous numerical features. I utilized **StandardScaler** exclusively on the five continuous variables: tenure, monthly charges, total charges, signal stability index, and outage history score. This normalization was necessary to ensure a legitimate evaluation, the StandardScaler was fitted only on the training set and then applied to the test set, preventing information from 'bleeding' into the unseen data. The variables with large numerical ranges (like total charges) did not disproportionately influence the model's weightings compared to smaller, yet critical, metrics (like signal stability index).

5.0.3 Data Leakage Prevention

A critical best practice was followed during the preprocessing pipeline to prevent target leakage. The churn variable was explicitly isolated and mapped to binary values before One-Hot Encoding and Feature Scaling. This ensures the features matrix (X) contains only predictive indicators. Furthermore, by **fitting the scaling parameters only on the training data**, I maintained strict separation of the test set, allowing for an unbiased evaluation of model performance.

6.0 Model Selection and Optimization

In order to analyze the interaction between behavioral data and the new geospatial risk factors, I undertook a rigorous model selection process. Churn in a hybrid environment is not merely a linear trend; it often occurs at "tipping points," such as when a customer’s signal stability drops below a specific threshold.

6.0.1 The Model Tournament (Cross-Validation)

I evaluated three models, namely Logistic Regression, Random Forest, and Gradient Boosting. For each of these models, I performed **GridSearch** to optimize hyperparameters and used 5-fold cross-validation to ensure stable results across different subsets of the data.

Model	Hyperparameters	Mean ROC-AUC (CV)	Rationale for Testing
Logistic Regression	penalty: elasticnet, C:10, L1_ratio : 1.0	0.8448	Winner for Interpretability and Performance
Gradient Boosting	n_estimators: 100, learning_rate: 0.1	0.8470	Highest raw score but used for comparison.
Random Forest	n_estimators: 200, max_depth: 10	0.8322	Lowest raw score but used for comparison.

Ultimately, Gradient Boosting performed marginally better (0.847); however, Logistic Regression (0.845) was selected as the final model. This choice balances high discriminative power with the transparency required to explain **why** specific geospatial factors—like signal stability—are triggering churn to non-technical stakeholders.

6.0.2 Metric Selection and Rationale

For the Hybrid Telco Churn Model, I primarily optimized for ROC-AUC, while closely monitoring Recall. Given that the goal is Geospatial Risk Analysis, it was essential to select a model capable of effectively ranking customers by their risk of churn, from highest to lowest. ROC-AUC measures how well the model separates the "Stay" and "Churn" classes across all possible thresholds, making it the primary metric. A score of 0.84 proves that the hybrid features (Signal and Outages) are powerful discriminators, allowing the business to prioritize neighborhoods that require infrastructure repairs. In telecom, the cost of losing a customer due to churn is much higher than the cost of a retention effort, such as sending a proactive text or fixing a tower.

The business motives include:

- False Negative (The "Miss"): Predicts that a customer will stay, but they leave. Cost: High (Lost revenue + cost to acquire a new customer).
- False Positive (The "False Alarm"): Predicts that a customer will leave, but they stay. Cost: Low (Small cost of a proactive service check or a discount).

By optimizing Recall, I ensure I "catch" as many potential leavers as possible, even if this results in a few false alarms. By selecting the **0.43 threshold**, this shift allows me to achieve that **85% Recall** (catching 85% of churners).

6.0.3 Model Optimization and Strategy

To transition from a general predictive model to a specialized **Hybrid Geospatial Risk** tool, I performed a two-step optimization focusing on hyperparameter tuning and decision threshold calibration.

6.0.4 Hyperparameter Tuning (Grid Search)

Using **GridSearchCV**, I exhaustively tested combinations of model settings, specifically, regularization strength (C) and the **ElasticNet mixing ratio (l1_ratio)**, to identify the configuration that best generalized the geospatial 'noise' injected into the dataset. This systematic refinement resulted in a final **Test ROC-AUC of 0.84**, confirming that the model remains highly

robust when encountering unseen data. Even without the complexity of deep trees, the Logistic Regression model proves that the hybrid features provide enough 'signal' to successfully distinguish between at-risk and loyal customers.

7.0 Results and Final Evaluation

After tuning, the final model was evaluated on the unseen Test Set, achieving an overall accuracy of 71% at the optimized threshold, which prioritizes the capture of 85% of at-risk customers.

7.0.1 Model Performance (ROC Curve)

To evaluate the overall diagnostic ability of our hybrid model, I utilized the ROC curve. This metric measures the model's performance across all possible classification thresholds, providing a comprehensive view of how well the system distinguishes between 'Stay' and 'Churn' customers.




Figure 6: Characteristic (ROC) Curve.

The graph shows a steep initial rise on the left, suggesting that the model achieved an AUC score of 0.84, which indicates excellent discriminative power. The curve rises steeply toward the top-left corner, proving that the model can maintain a high True Positive Rate (Recall) while keeping the False Positive Rate at a manageable ~35% even when targeting maximum recall.

For our Geospatial Risk Analysis, this high AUC score confirms that the combination of engineered signal stability metrics and traditional customer data has created a robust prediction

engine. We aren't just guessing; the model has a strong, mathematically proven ability to identify at-risk customers across different geographical regions.

7.0.2 Balancing Accuracy and Coverage (PR Curve)

While ROC-AUC measures overall model power, I utilized a **Precision-Recall (PR) Curve** to operationalize the model for business use. This chart allows us to identify the "sweet spot" where we catch the most churners (Recall) without overwhelming our budget with false alarms (Precision).




Figure 7: Balancing Accuracy and Coverage

The curve demonstrates that precision remains relatively stable (above 60%) until we attempt to capture more than 55% of churners. Beyond this point, the trade-off becomes steeper as the model begins to misclassify more stable customers. I intentionally selected a threshold that pushes Recall toward 85%. While this results in a Precision of ~47%, it represents a calculated **'Defensive Strategy.'** This prioritization ensures that the vast majority of at-risk revenue is identified, accepting a higher rate of 'false alarms' (approximately 1.1 false alarms for every 1 true positive) in exchange for **maximum risk coverage**. In a highly competitive telco market, catching nearly all churners (high recall) is more valuable than perfect accuracy (high precision). The cost of a proactive signal check or a retention 'bridge' discount is significantly lower than the **Churn Acquisition Cost (CAC)** required to replace a lost high-value subscriber.

7.0.3 Performance Diagnosis (Confusion Matrix)

To understand the real-world impact of our model, I analyzed the Confusion Matrix at the **Optimized Threshold (0.43)**. This visualization breaks down our successes (True Positives/Negatives) and our errors (False Positives/Negatives) to evaluate our risk coverage.




Figure 8: Confusion Matrix

The matrix shows:

- **True Positives (318):** We successfully identified 318 customers who were actually going to leave. These are our "Saves" high-value targets for immediate retention outreach.
- **False Negatives (56):** Only 56 churners were missed by the model. This low number is critical for our goal; we have minimized the "blind spots" in our churn detection.
- **False Positives (353):** While 353 customers were flagged as "at-risk" but stayed, this is a calculated strategic cost. In the Telecom industry, the cost of sending an automated retention offer or performing a remote signal check on a healthy customer is negligible compared to the \$500+ cost of losing a subscriber.
- **True Negatives (682):** Correct identification of loyal customers.

In summary, by shifting the decision threshold to 0.43, we have successfully captured **85% of all churners** (Recall), ensuring that the vast majority of at-risk revenue is identified for intervention before the customer disconnects.

7.0.4 Identification of Key Churn Drivers

To validate the effectiveness of our Hybrid & Geospatial approach, I extracted the coefficients from the optimized Logistic Regression model. This analysis identifies which variables, behavioral, financial, or technical carry the most weight in predicting a customer's decision to exit the service.




Figure 9: Top 10 Drivers of Churn.

The coefficient analysis reveals that Internet Service, fiber optic and monthly charges are the primary drivers of churn risk. Interestingly, while the Signal Stability Index does not appear in the global Top 10, our earlier EDA proves it is a 'High-Impact Localized Driver' in specific geospatial zones. The high ranking of contract two-year and tenure (acting as protective factors that decrease churn risk) confirms that behavioral 'anchors' are the strongest defense against churn. Furthermore, the presence of Streaming TV and Streaming Movies in the Top 10 suggests that 'Heavy Users,' those most dependent on high-bandwidth, stable connections, represent the most sensitive and high-value risk segment for the company. This validates our hybrid approach: while financial metrics identify *who* might leave, our technical indicators explain *why* high-value users in poor signal areas are more vulnerable.

8.0 Prescriptive Simulation (The ROI)

To transition from prediction to action, I conducted a "What-If" simulation to quantify the financial impact of infrastructure investment in the highest-risk zone.

8.0.1 The Intervention Strategy

I focused on customers in the identified "Trouble Zone" (Zip 98104) and simulated a technical intervention that improved signal stability from 0.45 to 0.80.

8.0.2 Quantifiable Results

Current Average Churn Risk (Zip 98104): 22.5%

Projected Average Risk (Post-Fix): 22.12%

Business Impact: A 1.69% reduction in regional churn.




Figure 10: Impact of Fixed Infrastructure.

8.0.3 ROI Conclusion

Our simulation highlights a reduction in churn risk from 22.5% to 22.12% in the high-risk zone, representing a 1.69% relative decrease in regional attrition through technical fixes alone. By utilizing a model with high interpretability, we can directly link the **1.69% relative churn reduction** in Zip 98104 to the specific mathematical weight of signal stability. This validates the Hybrid approach as a **prescriptive** tool for infrastructure investment. While this proves that infrastructure intervention is a powerful surgical tool, the remaining 22.12% risk suggests that signal quality is only part of the story. Therefore, we recommend a dual strategy: prioritizing

technical repairs to stabilize the connection while simultaneously offering contract incentives to anchor the customer behaviorally, ensuring long-term loyalty.

9.0 Conclusion and Actionable Recommendations

9.1 Actionable Recommendations

Based on the predictive power of the hybrid model, the following strategies are recommended to optimize customer retention:

- **Surgical Infrastructure Intervention:** Prioritize technical upgrades in the top 5% of geospatial churn hotspots (e.g., ZIP 98104). Since technical reliability acts as the "breaking point" for high-value Fiber Optic users, targeted enhancements in these zones will yield higher ROI than generic discounts.
- **Proactive Technical Triggers:** Integrate the **Signal Stability Index** and **Outage History Score** into real-time monitoring systems. Using these as early-warning triggers allows the technical team to resolve stability issues before they escalate into customer attrition.
- **Data-Driven Win-Back Campaigns:** Deploy targeted marketing to recently churned customers in "repaired" zones. By communicating specific infrastructure improvements made in their neighborhood, the company can provide a factual, high-incentive reason for customers to return.
- **Geospatial Loyalty Bridging:** Implement loyalty programs specifically for "developing infrastructure zones." Rewarding long-term patronage in these areas helps bridge the gap and anchor customers behaviorally while technical stability is being optimized.
- **Localized Feedback Loops:** Establish direct engagement channels for customers in low-signal areas. This qualitative data ensures that technical investments remain aligned with actual consumer experiences and expectations.

9.2 Conclusion

This Hybrid Telco Churn and Geospatial Risk Analysis demonstrates that while financial and behavioral metrics (tenure, contract type) identify who might leave, geospatial technical data explains why they leave. By optimizing a Logistic Regression model via **GridSearch**, this study

established a robust predictive engine, achieving 85% recall of at-risk customers. The prescriptive simulation proved that targeted infrastructure fixes can achieve a 1.69% relative reduction in regional churn. This confirms that infrastructure intervention is not a "blanket" solution, but a surgical tool to be deployed in high-value, low-stability zones. Ultimately, the evolution of churn management lies in this hybrid approach—leveraging data science to balance cost-effective technical repairs with behavioral incentives. Future expansions integrating social sentiment and real-time satisfaction scores will further refine this model, ensuring the company maintains a competitive edge in a shifting telecommunications landscape.